

# Introduction to the Design and Analysis of Randomized Experiments

## Class 4: Power Analysis

Jake Bowers

July 31, 2024

## Overview and Review

What is power?

Analytical calculations of power

Simulation-based power calculation

Power with covariate adjustment

Power for cluster randomization

Power Ingredients and Their Relationship

Pre-Registration of Analysis Plans

References

## Overview and Review

# Today

0. Quiz and Questions
1. Statistical Power (1 - false negative rate of tests)
2. Pre-registration of analysis plans

Note: You can download (and contribute to) course materials at  
[https://github.com/bowers-illinois-edu/short\\_course\\_experiments](https://github.com/bowers-illinois-edu/short_course_experiments)

Hay recursos en español aquí:

[https://egap.github.io/theory\\_and\\_practice\\_of\\_field\\_experiments\\_spanish/](https://egap.github.io/theory_and_practice_of_field_experiments_spanish/)

# Lingering Questions?

Questions arising?

# Quiz 1

What is wrong about these statements?

- “When we estimate the ATE using the difference of means, we have to assume that  $Y \sim \text{Normal}()$ .”
- “We use a linear model to estimate the ATE. You can't just subtract the mean outcome of the treatment group from the mean outcome of the control group.”

## Quiz 2

What is wrong about these statements?

- “When we write  $E_R[\widehat{ATE}] = ATE$  it means that our estimator of the ATE is unbiased. I looked up what ChatGPT says about “unbiased” and it says:

Here are some synonyms for “biased”:

1. Prejudiced
2. Partial
3. One-sided
4. Slanted
5. Skewed
6. Predisposed
7. Influenced
8. Tendentious
9. Unfair
10. Swayed

So, when we use an unbiased estimator we are not prejudiced or unfair. I also think that an unbiased estimator tells you the truth.”

## Quiz 2

What is wrong about these statements?

- “When you use a block randomized trial you have to use a fixed effects estimator, with fixed effects for the blocks.”
- “When you use a cluster randomized trial you have to use a multilevel model with random effects for the cluster.”
- “When you have a binary outcome, the only meaningful estimates come from a logit model.”



What is power?

# What is power?

We want to separate signal from noise.

- Power = probability of rejecting null hypothesis, given true effect  $\neq 0$ . We would like to have  $p \leq \alpha$  for the hypothesis of no effects when the truth is not zero.
- It is the ability to detect signal from noise (assuming there is a signal).
- Formally: power = (1 - Type II) error rate.
- Thus, power  $\in (0, 1)$ .
- How often should we see a  $p \leq \alpha$ ? Standard thresholds: 0.8 or 0.9 — “nearly always detect a signal when one exists”

# Starting point for power analysis

- Power analysis is something we do *before* we run a study.
  - ▶ Helps you figure out the sample you need to detect a given effect size.
  - ▶ Or helps you figure out a minimal detectable difference given a set sample size.
  - ▶ **May help you decide whether to run a study at all.** (A power analysis is part of answering the question, “Should we do this study?”)
- It is hard to learn from an under-powered null finding.
  - ▶ Was there an effect, but we were unable to detect it? or was there no effect? We can't say.
  - ▶ How should we interpret “The difference in proportion vaccinated between Message A and Control was .02 ( $p = .4$ ).”?

# Power

- Say there truly is a treatment effect and you run your experiment many times (hypothetically) with the same group of people. How often will you get a  $p \leq .05$ ?
- It depends:
  - ▶ How big is your treatment effect?
  - ▶ How many units are treated, measured?
  - ▶ How much noise is there in the measurement of your outcome?
- **We do not know the answers to all those questions in advance. So some guesswork required to answer this question.**

# Approaches to power calculation

- Analytical calculations of power
- Simulate: Repeat the experiment with guessed-at treatment effect sizes, outcome variability, and  $N$ .

# Power calculation tools

- Interactive
  - ▶ EGAP Power Calculator
  - ▶ rpsychologist
- R Packages
  - ▶ pwr
  - ▶ DeclareDesign, see also <https://declaredesign.org/>

## Analytical calculations of power

# Analytical calculations of power for hypotheses about no average treatment effects

- Formula:

$$\text{Power} = \Phi \left( \frac{|\tau|\sqrt{N}}{2\sigma} - \Phi^{-1}\left(1 - \frac{\alpha}{2}\right) \right)$$

- Components:

- ▶  $\Phi$ : standard normal CDF is monotonically increasing
- ▶  $\tau$ : the effect size
- ▶  $N$ : the sample size
- ▶  $\sigma$ : the standard deviation of the outcome
- ▶  $\alpha$ : the significance level (typically 0.05)

Why standard normal? (the CLT!)



## Example: Analytical calculations of power

*# Power for a study with 80 obserations and effect size of 0.25*

```
library(pwr)
pwr.t.test(
  n = 40, d = 0.25, sig.level = 0.05,
  power = NULL, type = c(
    "two.sample",
    "one.sample", "paired"
  )
)
```

```
##
##      Two-sample t test power calculation
##
##              n = 40
##              d = 0.25
##      sig.level = 0.05
##      power = 0.1971831
##      alternative = two.sided
##
## NOTE: n is number in *each* group
```

# Limitations to analytical power calculations

- Only derived for some estimands (ATE/ITT)
- Makes specific assumptions about the data-generating process (for example,  $N$  is large enough that the reference distribution for the test statistic is close to Normal)
- Difficult with more complex designs like block randomized designs with different probabilities of assignment in each block.

## Simulation-based power calculation

# Simulation-based power calculation steps

- Create dataset and simulate research design.
- Assumptions are necessary for simulation studies, but you make your own.
- For the DeclareDesign approach, see <https://declaredesign.org/>

# Steps

- Define the sample and the potential outcomes function.
- Define the treatment assignment procedure.
- Create data.
- Assign treatment, then estimate the effect.
- Do this many times.

# Examples

- Complete randomization
- With covariates
- With cluster randomization

## Example: Simulation-based power for complete randomization

```
# install.packages("randomizr")
library("randomizr")
library("estimatr")

## Y0 is fixed in most field experiments.
## So we only generate it once (here making it Normal parallels with
make_Y0 <- function(N) {
  rnorm(n = N)
}

repeat_experiment_and_test <- function(N, Y0, tau) {
  # N is size of experimental pool; Y0 is potential outcome to control
  # tau is effect size (here, a constant additive effect)
  Z <- complete_ra(N = N)
  Y1 <- Y0 + Z * tau
  Yobs <- Z * Y1 + (1 - Z) * Y0
  estimator <- lm_robust(Yobs ~ Z)
  pval <- estimator$p.value[2]
  return(pval)
}
```

## Example: Simulation-based power for complete randomization

```
power_sim <- function(N, tau, sims) {  
  Y0 <- make_Y0(N)  
  pvals <- replicate(  
    n = sims,  
    repeat_experiment_and_test(N = N, Y0 = Y0, tau = tau)  
  )  
  pow <- sum(pvals < .05) / sims  
  return(pow)  
}
```

```
set.seed(12345)  
## Notice simulation variability with sims=100  
power_sim(N = 80, tau = .25, sims = 1000)
```

```
## [1] 0.157
```

```
power_sim(N = 80, tau = .25, sims = 1000)
```

```
## [1] 0.237
```



## Example: Using DeclareDesign I

```
library(DeclareDesign)
library(tidyverse)
P0 <- declare_population(N, u0 = rnorm(N))
# declare Y(Z=1) and Y(Z=0)
O0 <- declare_potential_outcomes(Y_Z_0 = 5 + u0, Y_Z_1 = Y_Z_0 + tau)
# design is to assign m units to treatment
A0 <- declare_assignment(Z = conduct_ra(N = N, m = round(N / 2)))
# estimand is the average difference between Y(Z=1) and Y(Z=0)
estimand_ate <- declare_inquiry(ATE = mean(Y_Z_1 - Y_Z_0))
R0 <- declare_reveal(Y, Z)
design0_base <- P0 + A0 + O0 + R0

## For example with N=100 and tau=.25:
design0_N100_tau25 <- redesign(design0_base, N = 100, tau = .25)
dat0_N100_tau25 <- draw_data(design0_N100_tau25)
head(dat0_N100_tau25)
```

## Example: Using DeclareDesign II

```
##      ID      u0 Z      Y_Z_0      Y_Z_1      Y
## 1 001  1.2790709 0 6.279071 6.529071 6.279071
## 2 002  2.3056456 0 7.305646 7.555646 7.305646
## 3 003 -0.9621603 1 4.037840 4.287840 4.287840
## 4 004 -0.8225563 1 4.177444 4.427444 4.427444
## 5 005  0.1702612 1 5.170261 5.420261 5.420261
## 6 006 -0.1926055 1 4.807394 5.057394 5.057394
```

```
with(dat0_N100_tau25, mean(Y_Z_1 - Y_Z_0)) # true ATE
```

```
## [1] 0.25
```

```
with(dat0_N100_tau25, mean(Y[Z == 1]) - mean(Y[Z == 0])) # estimate
```

```
## [1] 0.2217877
```

```
lm_robust(Y ~ Z, data = dat0_N100_tau25)$coef # estimate
```

```
## (Intercept)      Z
##  5.0297228  0.2217877
```

## Example: Using DeclareDesign III

```
E0 <- declare_estimator(Y ~ Z,
  .method = lm_robust, label = "t test 1",
  inquiry = "ATE"
)

t_test <- function(data) {
  test <- with(data, t.test(x = Y[Z == 1], y = Y[Z == 0]))
  data.frame(statistic = test$statistic, p.value = test$p.value)
}

T0 <- declare_test(handler = label_test(t_test), label = "t test 2")

design0_plus_tests <- design0_base + E0 + T0

design0_N100_tau25_plus <- redesign(design0_plus_tests, N = 100, tau

## Only repeat the random assignment, not the creation of Y0. Ignore
names(design0_N100_tau25_plus)

## [1] "P0"          "A0"          "00"          "R0"          "t test 1" "t test 2"
```

## Example: Using DeclareDesign IV

```
design0_N100_tau25_sims <- simulate_design(design0_N100_tau25_plus,  
  sims = c(1, 100, 1, 1, 1, 1)  
) # only repeat the random assignment  
# design0_N100_tau25_sims has 200 rows (2 tests * 100 random assignments)  
# just look at the first 6 rows  
head(design0_N100_tau25_sims)
```

```
##           design      N   tau sim_ID estimator term  estimate  
## 1 design0_N100_tau25_plus 100 0.25      1  t test 1      Z 0.2040486  
## 2 design0_N100_tau25_plus 100 0.25      1  t test 2 <NA>          NA  
## 3 design0_N100_tau25_plus 100 0.25      2  t test 1      Z 0.2000803  
## 4 design0_N100_tau25_plus 100 0.25      2  t test 2 <NA>          NA  
## 5 design0_N100_tau25_plus 100 0.25      3  t test 1      Z 0.3892886  
## 6 design0_N100_tau25_plus 100 0.25      3  t test 2 <NA>          NA  
## statistic      p.value      conf.low conf.high df outcome inquiry s  
## 1  1.040959 0.30045567 -0.184946209 0.5930435 98      Y      ATE  
## 2  1.040959 0.30054813          NA          NA NA      <NA>      <NA>  
## 3  1.020767 0.30987837 -0.188894886 0.5890555 98      Y      ATE  
## 4  1.020767 0.30987957          NA          NA NA      <NA>      <NA>
```

## Example: Using DeclareDesign V

```
## 5  1.990541 0.04931644 0.001187836 0.7773893 98      Y      ATE
## 6  1.990541 0.04938403      NA      NA NA    <NA>    <NA>
##    step_2_draw
## 1          1
## 2          1
## 3          2
## 4          2
## 5          3
## 6          3
```

## Power with complete randomization

In 26% of experiments, when the truth is  $.25sds$ , and  $N = 100$ , we get  $p < .05$ .

```
## # A tibble: 2 x 2
##   estimator    pow
##   <chr>      <dbl>
## 1 t test 1    0.26
## 2 t test 2    0.26
```

## Power with covariate adjustment

# Covariate adjustment and power

- Covariate/Covariance adjustment can improve power because it mops up variation in the outcome variable.
  - ▶ If prognostic (predictive of the outcome), covariate adjustment can reduce variance dramatically. Lower outcome variance means higher power.
  - ▶ If non-prognostic, power gains are minimal.
- All covariates must be pre-treatment. Do not drop observations on account of missingness.
  - ▶ See the [Theory and Practice module on threats to internal validity](#) and the [10 things to know about covariate adjustment](#).
- Freedman (2008) pointed out that covariance-adjusted estimators of the ATE are biased. Lin (2013) shows that bias decreases with  $N$ .



# Blocking

- Blocking: randomly assign treatment within blocks
  - ▶ “Ex-ante” covariate adjustment
  - ▶ Higher precision/efficiency implies more power
  - ▶ Reduce “conditional bias”: association between treatment assignment and potential outcomes
  - ▶ Benefits of blocking over covariate adjustment clearest in small experiments

## Example: Simulation-based power with a covariate I

```
## Y0 is fixed in most field experiments. So we only generate it once
make_Y0_cov <- function(N) {
  u0 <- rnorm(n = N)
  x <- rpois(n = N, lambda = 2)
  Y0 <- .5 * sd(u0) * x + u0
  return(data.frame(Y0 = Y0, x = x))
}
## X is moderately predictive of Y0.
test_dat <- make_Y0_cov(100)
test_lm <- lm_robust(Y0 ~ x, data = test_dat)
summary(test_lm)
```

```
##
## Call:
## lm_robust(formula = Y0 ~ x, data = test_dat)
##
## Standard error type: HC2
##
## Coefficients:
```

## Example: Simulation-based power with a covariate II

```
##           Estimate Std. Error t value Pr(>|t|) CI Lower CI Upper
## (Intercept) -0.1223    0.17470 -0.7001 4.856e-01  -0.469    0.224
## x           0.5387    0.07341  7.3385 6.354e-11   0.393    0.684
##
## Multiple R-squared:  0.3716 ,    Adjusted R-squared:  0.3652
## F-statistic: 53.85 on 1 and 98 DF,  p-value: 6.354e-11
```

## Example: Simulation-based power with a covariate III

```
## now set up the simulation
repeat_experiment_and_test_cov <- function(N, tau, Y0, x) {
  Z <- complete_ra(N = N)
  Y1 <- Y0 + Z * tau
  Yobs <- Z * Y1 + (1 - Z) * Y0
  estimator <- lm_robust(Yobs ~ Z + x, data = data.frame(Y0, Z, x))
  pval <- estimator$p.value[2]
  return(pval)
}

## create the data once, randomly assign treatment sims times
## report what proportion return p-value < 0.05
power_sim_cov <- function(N, tau, sims) {
  dat <- make_Y0_cov(N)
  pvals <- replicate(n = sims, repeat_experiment_and_test_cov(
    N = N,
    tau = tau, Y0 = dat$Y0, x = dat$x
  ))
  pow <- sum(pvals < .05) / sims
  return(pow)
}
```

## Example: Simulation-based power with a covariate I

Doing it twice to be clear that there is variability from simulation to simulation.

```
set.seed(12345)
power_sim_cov(N = 80, tau = .25, sims = 100)
```

```
## [1] 0.13
```

```
power_sim_cov(N = 80, tau = .25, sims = 100)
```

```
## [1] 0.19
```

## Power for cluster randomization

# Power and clustered designs

- Given a fixed  $N$ , a clustered design is often less powered than a non-clustered design.
  - ▶ The difference is often substantial.
- We have to estimate variance correctly:
  - ▶ Clustering standard errors (the usual)
  - ▶ Randomization inference
- To increase power:
  - ▶ Better to increase number of clusters than number of units per cluster if treatment is at level of cluster.
  - ▶ How much clusters reduce power depends critically on the intra-cluster correlation (the ratio of variance within clusters to total variance).

## Example: Simulation-based power for cluster randomization

|

```
## Y0 is fixed in most field experiments. So we only generate it once
make_Y0_clus <- function(n_indivs, n_clus) {
  # n_indivs is number of people per cluster
  # n_clus is number of clusters
  clus_id <- gl(n_clus, n_indivs)
  N <- n_clus * n_indivs
  u0 <- fabricatr::draw_normal_icc(N = N, clusters = clus_id, ICC = .5)
  Y0 <- u0
  return(data.frame(Y0 = Y0, clus_id = clus_id))
}
```

```
test_dat <- make_Y0_clus(n_indivs = 10, n_clus = 100)
# confirm that this produces data with 10 in each of 100 clusters
table(test_dat$clus_id)
```



# Example: Simulation-based power for cluster randomization

## II

```
##
##   1    2    3    4    5    6    7    8    9   10   11   12   13   14   15   16   17
##  10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10
##  21   22   23   24   25   26   27   28   29   30   31   32   33   34   35   36   37
##  10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10
##  41   42   43   44   45   46   47   48   49   50   51   52   53   54   55   56   57
##  10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10
##  61   62   63   64   65   66   67   68   69   70   71   72   73   74   75   76   77
##  10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10
##  81   82   83   84   85   86   87   88   89   90   91   92   93   94   95   96   97
##  10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10   10
```

```
# confirm ICC
ICC::ICCbare(y = Y0, x = clus_id, data = test_dat)
```

```
## [1] 0.09654799
```

## Example: Simulation-based power for cluster randomization

### III

```
repeat_experiment_and_test_clus <- function(N, tau, Y0, clus_id) {  
  # here we randomize Z at the cluster level  
  Z <- cluster_ra(clusters = clus_id)  
  Y1 <- Y0 + Z * tau  
  Yobs <- Z * Y1 + (1 - Z) * Y0  
  estimator <- lm_robust(Yobs ~ Z,  
    clusters = clus_id,  
    data = data.frame(Y0, Z, clus_id), se_type = "CR2"  
  )  
  pval <- estimator$p.value[2]  
  return(pval)  
}  
  
power_sim_clus <- function(n_indivs, n_clus, tau, sims) {  
  dat <- make_Y0_clus(n_indivs, n_clus)  
  N <- n_indivs * n_clus  
  # randomize treatment sims times  
  pvals <- replicate(  
    n = sims,  
    repeat_experiment_and_test_clus(  
      n_indivs, n_clus, tau, dat$clus_id
```

## Example: Simulation-based power for cluster randomization (DeclareDesign) I

```
P1 <- declare_population(  
  N = n_clus * n_indivs,  
  clusters = gl(n_clus, n_indivs),  
  u0 = draw_normal_icc(N = N, clusters = clusters, ICC = .2)  
)  
O1 <- declare_potential_outcomes(Y_Z_0 = 5 + u0, Y_Z_1 = Y_Z_0 + tau)  
A1 <- declare_assignment(Z = conduct_ra(N = N, clusters = clusters))  
estimand_ate <- declare_inquiry(ATE = mean(Y_Z_1 - Y_Z_0))  
R1 <- declare_reveal(Y, Z)  
design1_base <- P1 + A1 + O1 + R1 + estimand_ate  
  
## For example:  
design1_test <- redesign(design1_base, n_clus = 10, n_indivs = 100, t  
test_d1 <- draw_data(design1_test)  
# confirm all individuals in a cluster have the same treatment assign  
with(test_d1, table(Z, clusters))
```

## Example: Simulation-based power for cluster randomization (DeclareDesign) II

```
##      clusters
## Z      1    2    3    4    5    6    7    8    9   10
##   0 100    0 100 100 100    0    0 100    0    0
##   1    0 100    0    0    0 100 100    0 100 100
```

## Example: Simulation-based power for cluster randomization (DeclareDesign) III

```
# four estimators, differ in se_type:
E1a <- declare_estimator(Y ~ Z,
  .method = lm_robust, clusters = clusters,
  se_type = "CR2", label = "CR2 cluster t test",
  inquiry = "ATE"
)
E1b <- declare_estimator(Y ~ Z,
  .method = lm_robust, clusters = clusters,
  se_type = "CR0", label = "CR0 cluster t test",
  inquiry = "ATE"
)
E1c <- declare_estimator(Y ~ Z,
  .method = lm_robust, clusters = clusters,
  se_type = "stata", label = "stata RCSE t test",
  inquiry = "ATE"
)
E1d <- declare_estimator(Y ~ Z,
  .method = lm_robust,
  se_type = "classical", label = "plain IID OLS t test",
  inquiry = "ATE"
```

## Example: Simulation-based power for cluster randomization (DeclareDesign) I

```
## Only repeat the random assignment, not the creation of Y0. Ignore  
## We would want more simulations in practice.  
set.seed(12355)  
design1_sims <- simulate_design(design1_plus_tosim,  
  sims = c(1, 1000, rep(1, length(design1_plus_tosim) - 2))  
)
```

# Simulation-based power for cluster randomizatio

Notice: high power but low coverage for plain OLS ("coverage" of a confidence interval is the same as false positive rate of a hypothesis test.)

```
design1_sims %>%  
  group_by(estimator) %>%  
  summarize(  
    pow = mean(p.value < .05),  
    coverage = mean(estimand <= conf.high & estimand >= conf.low),  
    .groups = "drop"  
  )
```

```
## # A tibble: 4 x 3  
##   estimator      pow coverage  
##   <chr>      <dbl>   <dbl>  
## 1 CR0 cluster t test  0.155   0.911  
## 2 CR2 cluster t test  0.105   0.936  
## 3 plain IID OLS t test 0.723   0.333  
## 4 stata RCSE t test   0.131   0.918
```

## Example: Simulation-based power for cluster randomization (DeclareDesign) I

```
## This may be simpler than the above:
library(DesignLibrary)
d1 <- block_cluster_two_arm_designer(
  N_blocks = 1,
  N_clusters_in_block = 10,
  N_i_in_cluster = 100,
  sd_block = 0,
  sd_cluster = .3,
  ate = .25
)
d1_plus <- d1 + E1b + E1c + E1d
d1_sims <- simulate_design(d1_plus, sims = c(1, 1, 1000, 1, 1, 1, 1,
```



## Example: Simulation-based power for cluster randomization (DeclareDesign) II

```
d1_sims %>%  
  group_by(estimator) %>%  
  summarize(  
    pow = mean(p.value < .05),  
    coverage = mean(estimand <= conf.high & estimand >= conf.low),  
    .groups = "drop"  
  )
```

```
## # A tibble: 4 x 3
```

##	estimator	pow	coverage
##	<chr>	<dbl>	<dbl>
## 1	CRO cluster t test	0.209	0.914
## 2	estimator	0.143	0.941
## 3	plain IID OLS t test	0.707	0.464
## 4	stata RCSE t test	0.194	0.925

## Power Ingredients and Their Relationship


# Comparative Statics

- Power is:
  - ▶ Increasing in  $N$
  - ▶ Increasing in  $|\tau|$
  - ▶ Decreasing in  $\sigma$

# Power by sample size I

```
some_ns <- seq(10, 800, by = 10)
pow_by_n <- sapply(some_ns, function(then) {
  pwr.t.test(n = then, d = 0.25, sig.level = 0.05)$power
})
plot(some_ns, pow_by_n,
     xlab = "Sample Size",
     ylab = "Power"
)
abline(h = .8)
```

## Power by sample size II




```
## See https://cran.r-project.org/web/packages/pwr/vignettes/pwr-vignettes  
## for fancier plots  
## ptest <- pwr.t.test(n = NULL, d = 0.25, sig.level = 0.05, power = 0.8)  
## plot(ptest)
```

## Power by treatment effect size $I$

```
some_taus <- seq(0, 1, by = .05)
pow_by_tau <- sapply(some_taus, function(theta) {
  pwr.t.test(n = 200, d = theta, sig.level = 0.05)$power
})
plot(some_taus, pow_by_tau,
     xlab = "Average Treatment Effect (Standardized)",
     ylab = "Power"
)
abline(h = .8)
```

# Power by treatment effect size II



# EGAP Power Calculator

- The calculator at: <https://egap.shinyapps.io/power-app/>
- For cluster randomization designs, try adjusting:
  - ▶ Number of clusters
  - ▶ Number of units per clusters
  - ▶ Intra-cluster correlation
  - ▶ Treatment effect



# Comments

- Know your outcome variable: what drives its variation regardless of treatment?
- What effects can you realistically expect from your treatment? What effects are substantively too small? (It may not be worth running *this experiment* at *this moment* if you cannot imagine detecting an effect at least this large.)
- What is the plausible range of variation of the outcome variable?
  - ▶ A design with limited possible movement in the outcome variable may not be well-powered.
  - ▶ See [10 Things Your Null Result Might Mean](#) for discussion of the various reasons a given experiment might have produced a  $p > .05$  (if you are using  $\alpha = .05$  as a rejection criteria or if you have some substantively small  $\hat{\tau}$  that might not be reached).

## Conclusion: How to improve your power

1. Increase the  $N$ 
  - ▶ If clustered, increase the number of clusters if at all possible
2. Strengthen the treatment effect ( $\tau_i$  and/or  $\bar{\tau}$ ) (What might this mean in your study?)
3. Improve precision by removing extraneous noise from the outcome
  - ▶ Covariate adjustment
  - ▶ Blocking
4. Better measurement of the outcome variable (for example, using indicies to reduce noise in the outcome variable)

# Pre-Registration of Analysis Plans

# Bias in published research against null results I

A good design executed well will produce credible research, which might be a null result. We want credible and actionable null results.

- Manuscripts with null results are never submitted for review or put away in a “file drawer” after rejections.
- We face incentives to change your specifications, measurements, or even hypotheses to get a statistically significant result ( $p$ -hacking and HARKing) to improve chances of publication.
  - ▶  $p$ -hacking: Trying many hypothesis tests increases the chances of a false positive result. (See Class 2 on Hypothesis Testing Slides)
  - ▶ HARKing: (Hypotheses written After Results Known) Pretending that you are **testing** hypotheses when you are not (you are generating them — which is fine but a separate activity).
- Even people not facing these incentives make many decisions when they analyze data: handling missing values and duplicate observations, creating scales, etc. Different seemingly small decisions can produce substantively meaningful differences in published results.

## Bias in published research against null results II

- Overall result: reduced credibility for individual pieces of research and (rightly) reduced confidence in whether we actually know what we claim to know.

# Pre-registration of analysis plans and research designs I

- Pre-registration is the filing of your research design and hypotheses with a publicly-accessible repository with a credible date stamp. EGAP hosts one that you can use for free. Same with OSF. (EGAP registry lives within the larger OSF registry). See also the [OES Analysis Plans](#).
- **Pre-registration does not preclude later exploratory analyses that were not stated in advance.** You just have to clearly distinguish between the two.

# Pre-registration of analysis plans and research designs II

- Even if you will be submitting a paper with results to an academic journal or you are primarily interested in a final report with findings for a policy audience, there are important advantages to you and to other researchers from pre-registering your research.
  - ▶ You can learn about other research, completed and in progress; others can learn about yours. We can learn about studies that produced null results.
  - ▶ It forces you to state your hypotheses and plan of analysis in advance of seeing the results, which limits  $p$ -hacking and HARKing.

## References



# References I



Freedman, David A. (2008). “On regression adjustments to experimental data”. In: *Advances in Applied Mathematics* 40.2, pp. 180–193.



Lin, Winston (Mar. 2013). “Agnostic notes on regression adjustments to experimental data: Reexamining Freedman’s critique”. en. In: *The Annals of Applied Statistics* 7.1, pp. 295–318. ISSN: 1932-6157. (Visited on 10/07/2016).